

Computing Quasi-Normal Modes of Schwarzschild Black Holes with Physics-Informed Neural Networks

Supervisor Progress Presentation

Jonathan Chung

University of Cambridge

25 February 2026

Outline

- 1 Black Hole Ringdown & the Master Equation
- 2 Methodology: PINNs & Finite Differences
- 3 Results

What Are Quasi-Normal Modes?

- A perturbed Schwarzschild black hole emits gravitational radiation as it settles down — the **ringdown**.
- The ringdown consists of **quasi-normal modes** (QNMs): damped sinusoids

$$\Phi(t) \propto e^{-t/\tau} \cos(\omega t)$$

- The frequency ω and decay time τ depend **only on the black hole mass M** (for Schwarzschild).
- **Goal:** Solve the perturbation equations numerically to extract ω and τ .

From Einstein's Equations to a Master Equation

Start with the Schwarzschild metric and add a small perturbation:

$$g_{\mu\nu} = g_{\mu\nu}^0(\text{Schwarzschild}) + h_{\mu\nu}(\text{small})$$

Key simplifications:

- ① Linearise Einstein's equations in $h_{\mu\nu}$.
- ② Decompose $h_{\mu\nu}$ in **tensor spherical harmonics** — angular dependence separates.
- ③ Construct a gauge-invariant **master function** Φ from the metric perturbation components.

Result: a single 1+1D wave equation for each angular mode ℓ :

$$\boxed{-\frac{\partial^2 \Phi}{\partial t^2} + \frac{\partial^2 \Phi}{\partial x^2} - V(r) \Phi = 0}$$

where $x = r + 2M \ln\left(\frac{r}{2M} - 1\right)$ is the **tortoise coordinate**.

The Potentials

The potential $V(r)$ depends on the **parity** of the perturbation:

Odd-parity (axial) — Regge–Wheeler:

$$V_{\text{RW}} = \left(1 - \frac{2M}{r}\right) \left[\frac{\ell(\ell+1)}{r^2} - \frac{6M}{r^3} \right]$$

Even-parity (polar) — Zerilli:

$$V_{\text{Z}} = \left(1 - \frac{2M}{r}\right) \frac{2n^2(n+1)r^3 + 6n^2Mr^2 + 18nM^2r + 18M^3}{r^3(nr+3M)^2}$$

where $2n = (\ell - 1)(\ell + 2)$.

- Both potentials peak near $x \sim 1M$ and vanish at the horizon and infinity.
- A transformation connects them (Chandrasekhar) \Rightarrow **same QNMs**.
- We focus on the **Zerilli potential** with $\ell = 2$ (dominant mode).

Boundary & Initial Conditions

Boundary conditions (Sommerfeld — radiation conditions):

$$\text{Horizon } (x \rightarrow -\infty) : (\partial_t - \partial_x) \Phi = 0 \quad (\text{ingoing})$$

$$\text{Infinity } (x \rightarrow +\infty) : (\partial_t + \partial_x) \Phi = 0 \quad (\text{outgoing})$$

Initial conditions — outgoing Gaussian pulse:

$$\Phi(x, 0) = \exp\left[-\frac{(x - 4M)^2}{(5M)^2}\right]$$

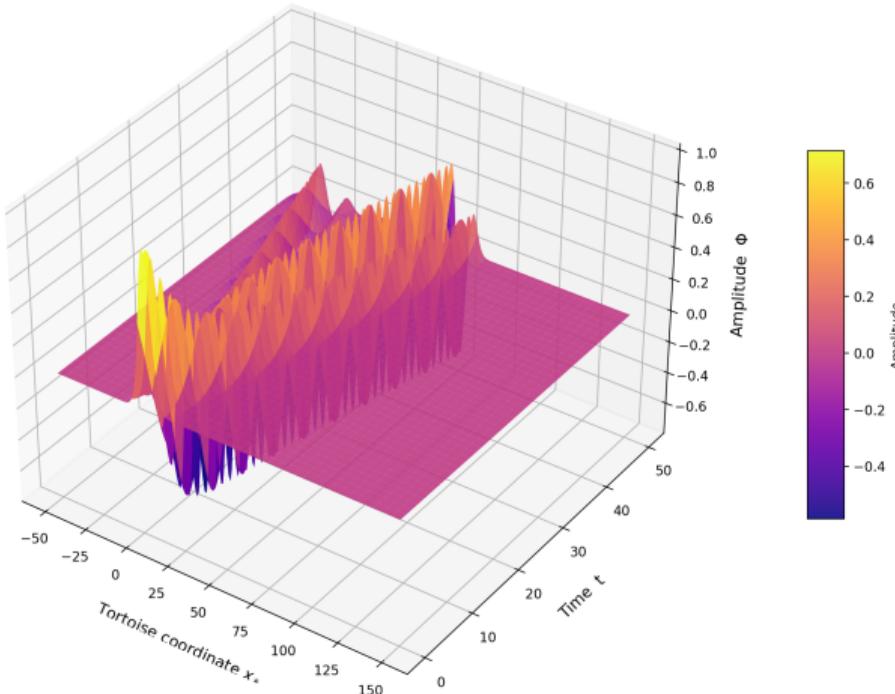
$$\partial_t \Phi(x, 0) = -\partial_x \Phi(x, 0) \quad (\text{purely right-moving})$$

Domain: $x/M \in [-50, 150]$, $t/M \in [0, 50]$.

The pulse scatters off the potential barrier \Rightarrow ringdown at late times.

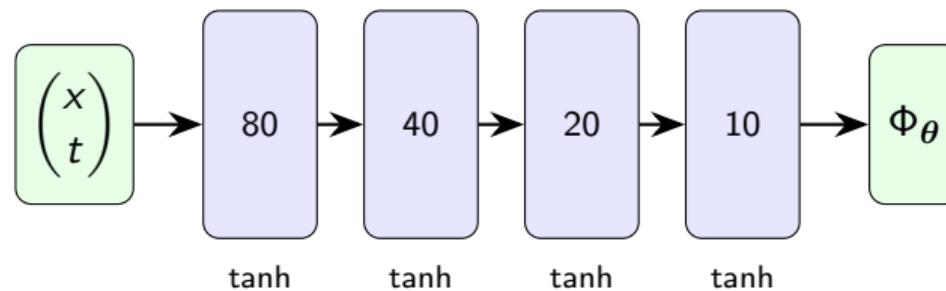
Ringdown Waveform — 3D Visualisation

3D Waveform of Schwarzschild Black Hole Ringdown (FD Reference)



Physics-Informed Neural Networks — Core Idea

Instead of a grid, train a neural network $\Phi_\theta(x, t)$ to satisfy the PDE, BCs, and ICs simultaneously.



- 4 hidden layers, tanh activation, Glorot uniform initialisation.
- Output transform: $\Phi_\theta = A \tanh(\text{raw output})$ bounds the solution to $[-1, 1]$.
- Derivatives via **automatic differentiation** (exact, no discretisation error).
- **4,521 trainable parameters.**

The Loss Function

Train by minimising a weighted sum of 7 terms:

$$\mathcal{L} = \lambda \cdot \left[\underbrace{\mathcal{L}_r, \mathcal{L}_{r_x}, \mathcal{L}_{r_t}}_{\text{PDE residual}}, \underbrace{\mathcal{L}_{ic}, \mathcal{L}_{iv}}_{\text{initial condns.}}, \underbrace{\mathcal{L}_{bl}, \mathcal{L}_{br}}_{\text{boundary condns.}} \right]$$

Term	Penalises	Points	λ
\mathcal{L}_r	$\Phi_{tt} - \Phi_{xx} + V\Phi \neq 0$	$N_r = 32,000$	100
$\mathcal{L}_{r_x}, \mathcal{L}_{r_t}$	Gradients of residual (gPINN)	same	100, 100
\mathcal{L}_{ic}	Initial profile mismatch	$N_i = 800$	1
\mathcal{L}_{iv}	Initial velocity mismatch	same	100
$\mathcal{L}_{bl}, \mathcal{L}_{br}$	Sommerfeld BC violations	$N_b = 400$	1, 1

- **Phase problem:** $\Phi(x, t + \alpha)$ is also a solution \Rightarrow weight λ_{iv} heavily to lock the correct phase.
- Total: **33,600 collocation points** (vs. 500,000 grid points for FD).

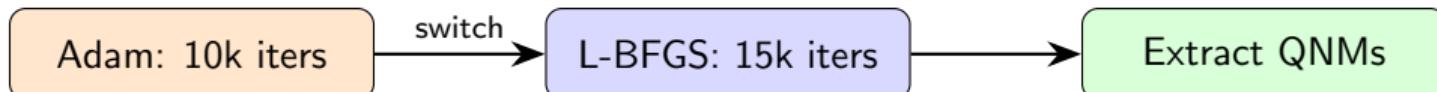
Two-Phase Training

Phase 1: Adam (10,000 iters)

- Stochastic gradient descent with adaptive learning rate ($\text{lr} = 10^{-3}$).
- Good at rough, global exploration.
- Uniform point resampling every 100 iters.

Phase 2: L-BFGS (15,000 iters)

- Quasi-Newton: approximates inverse Hessian from gradient history.
- Fast, precise convergence near minima.
- Sensitive to point changes \Rightarrow only uniform resampling here.



Finite Difference Baseline

For comparison, we solve the same PDE with standard numerical methods:

- Uniform mesh: $N_x = 1,000$ points, $\Delta x = 0.2M$.
- Method of lines + 4th-order Runge–Kutta, $\Delta t = 0.1M$.
- 500 time steps $\Rightarrow N_F = 500,000$ grid points.

The FD solution serves as **ground truth** for evaluating the PINN.

Key comparison:

	FD	PINN
Representation	Grid values $\Phi_{i,j}$	Neural network $\Phi_\theta(x, t)$
Derivatives	FD stencils (approx.)	Autograd (exact)
Points	500,000	33,600
Time-stepping	Explicit RK4	None (global solve)

Extracting Quasi-Normal Modes

Sample $\Phi(x_q, t)$ at observation point $x_q = 10M$, restrict to late times.

Method 1 — FFT + Envelope:

- ω : FFT of $\Phi(x_q, t) \Rightarrow$ peak frequency.
- τ : Linear fit to $\ln |\text{envelope peaks}|$ vs. $t \Rightarrow$ slope $= -1/\tau$.

Method 2 — Direct curve fit:

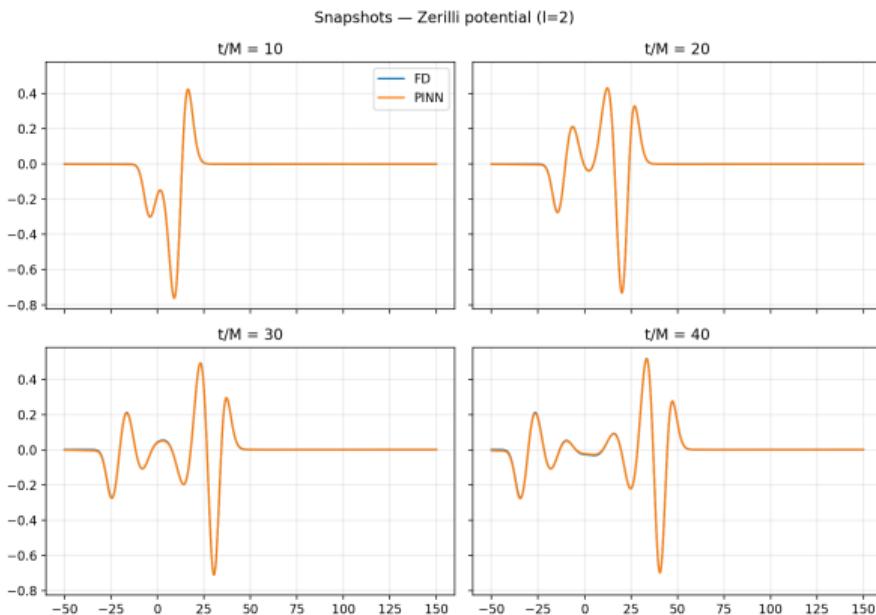
- Fit $\Phi \approx A e^{-t/\tau} \cos(\omega t + \phi)$ directly via nonlinear least squares.

Compare extracted (ω, τ) against theoretical values (Leaver 1985):

ℓ	ωM	τ/M
2	0.3737	11.241

Paper Reproduction: Baseline Results

Reproduced the setup of Patel, Aykutalp & Laguna (2024) with our corrected outgoing initial velocity profile.



Zerilli $\ell = 2$ QNM Extraction

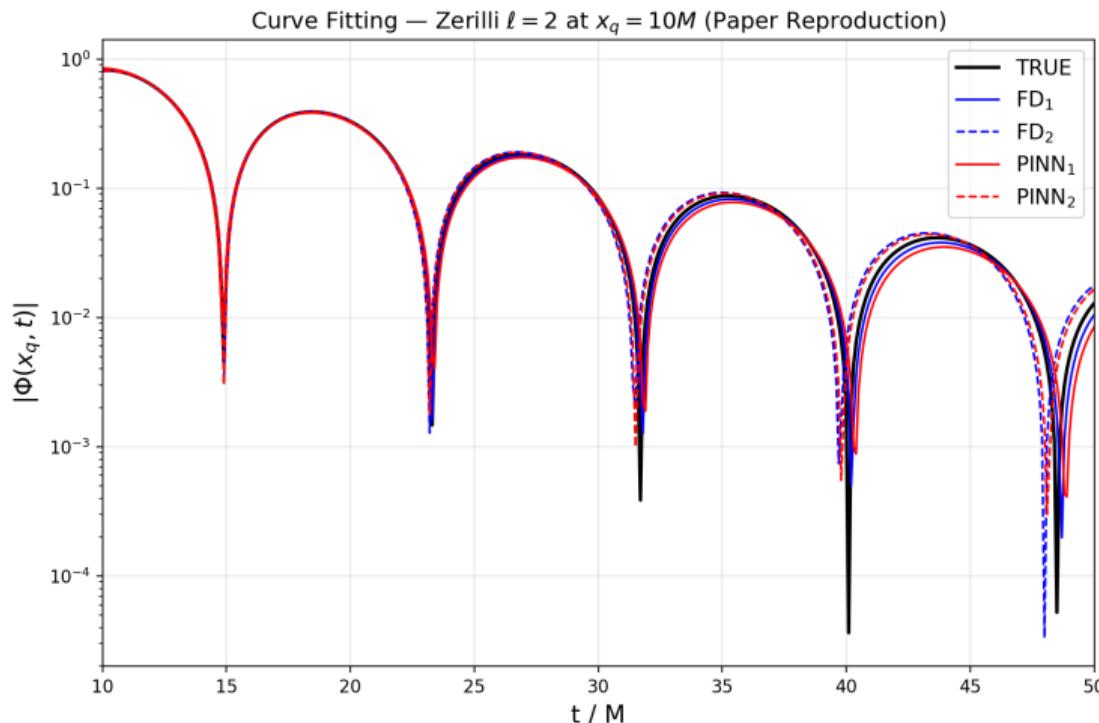
	τ/M	ωM
TRUE	11.241	0.3737
FD₁	10.929 (2.78)	0.3716 (0.56)
FD₂	11.465 (2.00)	0.3798 (1.63)
PINN₁	10.671 (5.07)	0.3697 (1.08)
PINN₂	11.410 (1.50)	0.3784 (1.25)

Parentheses: % error vs. Leaver 1985.

Subscript 1 = FFT + envelope.

Subscript 2 = direct curve fit.

Curve Fitting — Paper Reproduction



$\log |\Phi|$ vs. t/M at $x_q = 10M$. All five reconstructed damped cosines $A e^{-t/\tau} \cos(\omega t + \phi)$ overlap closely — both FD and PINN extract QNMs within $\sim 1\text{--}4\%$ of the Leaver (1985) values.

QNM Accuracy: Our Results vs. the Paper

Zerilli $\ell = 2$ — τ/M and ωM percentage errors

	Paper (Patel et al.)		Ours	
	τ/M (% err)	ωM (% err)	τ/M (% err)	ωM (% err)
FD ₁	10.804 (3.89)	0.370 (0.90)	10.929 (2.78)	0.372 (0.56)
FD ₂	11.073 (1.49)	0.378 (1.26)	11.465 (2.00)	0.380 (1.63)
PINN ₁	10.089 (10.25)	0.375 (0.29)	10.671 (5.07)	0.370 (1.08)
PINN ₂	9.620 (14.42)	0.376 (0.52)	11.410 (1.50)	0.378 (1.25)

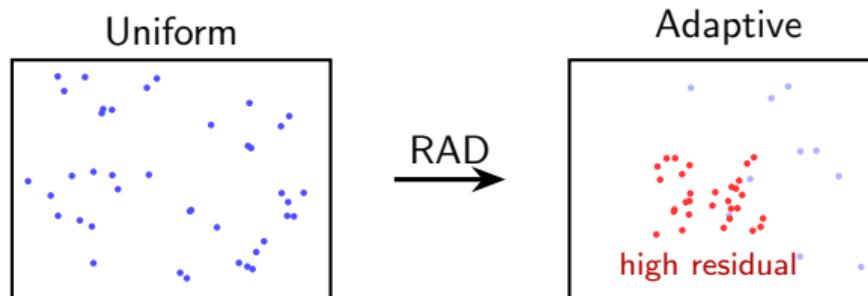
- **Major improvement:** PINN₂ decay time error drops from 14.42% → 1.50% (**9.6**× more accurate).
- PINN₁ τ error also improves: 10.25% → 5.07% (2.0×).
- Likely cause: our corrected outgoing initial velocity $\partial_t \Phi = -\partial_x \Phi$ (Sommerfeld condition).

Our Improvement: Residual-Adaptive Distribution (RAD)

Problem: Uniform sampling wastes points in empty regions where $\Phi \approx 0$.

RAD (Wu et al. 2023): Every P training steps during Adam—

- ① Generate 100,000 random candidate points.
- ② Evaluate PDE residual $|r_i|$ at each candidate.
- ③ Compute sampling probability: $p_i \propto |r_i|^k + c$.
- ④ Sample 32,000 points from candidates \Rightarrow replace all domain points.

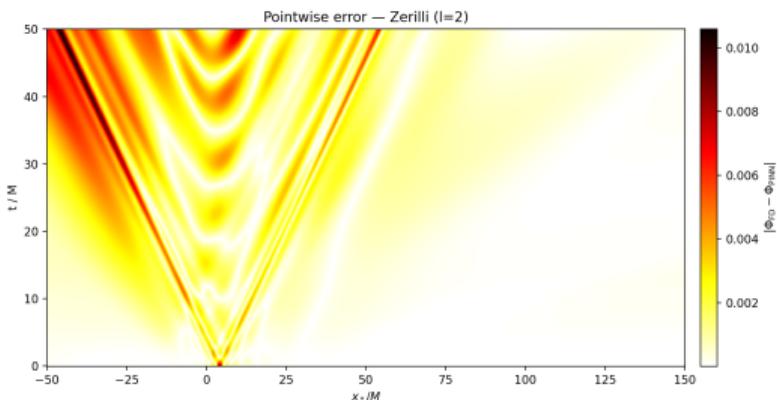


RAD Tuning Experiments

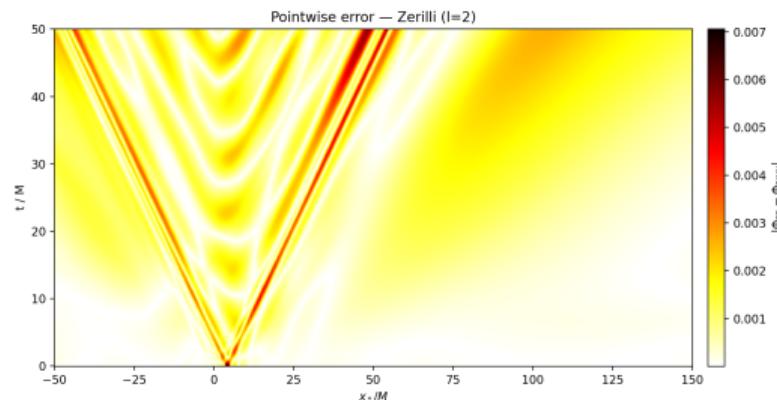
Configuration	k	c	Final Loss	RMSD	RL2
Uniform (paper repro)	—	—	1.18×10^{-6}	0.00183	1.30%
RAD baseline	1	1	7.40×10^{-7}	0.00115	0.82%
RAD aggressive	2	0.1	9.07×10^{-7}	0.00095	0.67%
RAD $k=2$, $P=500$	2	0.1	1.06×10^{-6}	0.00105	0.75%
RAD + anchor (20%)	1	1	8.20×10^{-7}	0.00141	1.00%

- RAD reduces RMSD by **37%** over uniform sampling.
- Aggressive RAD ($k=2$) achieves the **best RMSD** (0.00095).
- Anchor retention has the **second-lowest loss** (8.20×10^{-7}) but the **highest RMSD** among RAD runs — suggests over-fitting to anchor points at the expense of global accuracy.

Spatial Error Distribution



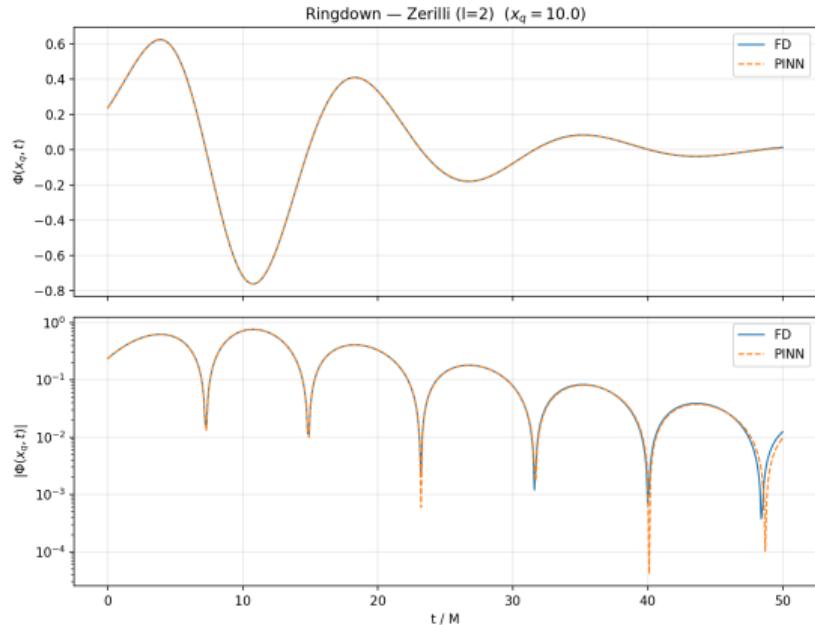
Uniform sampling



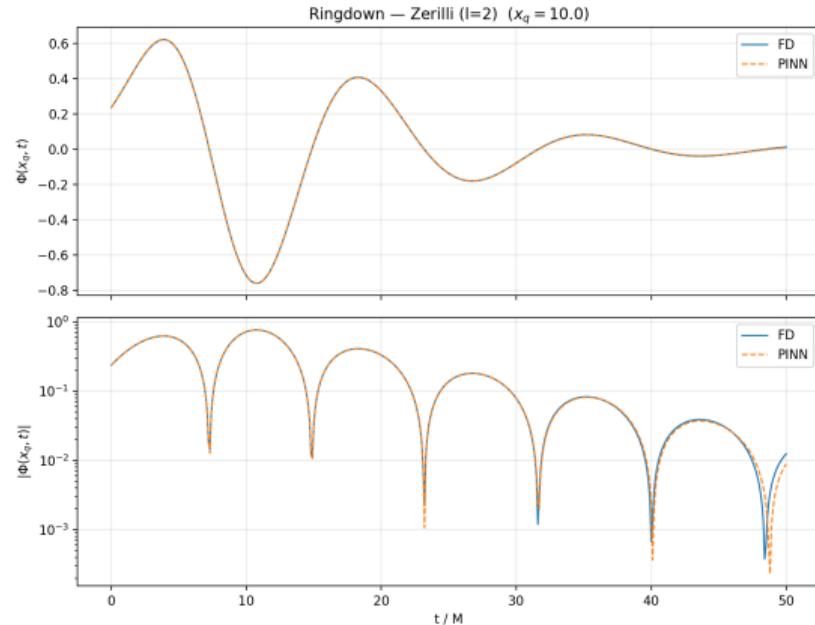
RAD adaptive sampling

RAD reduces the error concentration at the **wavefront leading edge**.

Ringdown Waveform at $x_q = 10M$



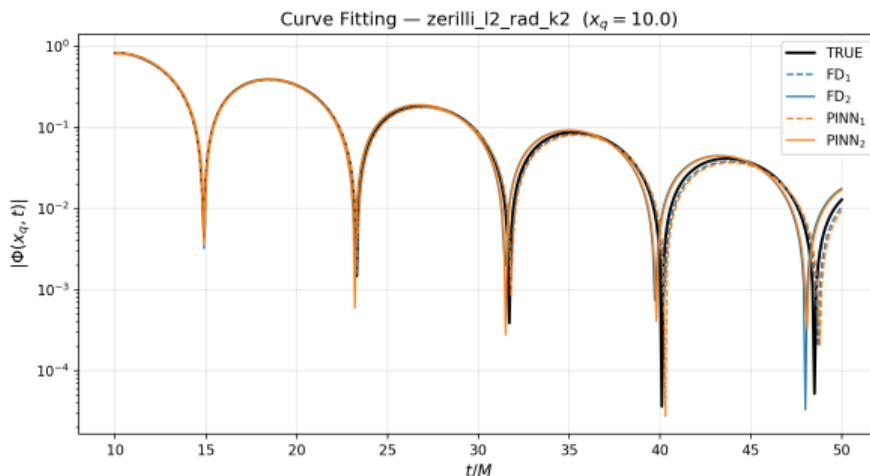
RAD baseline: PINN vs. FD (raw waveforms)



RAD $k = 2$: PINN vs. FD (raw waveforms)

Top: linear scale; Bottom: $\log |\Phi|$. Both variants closely match the FD reference.

Curve Fitting — RAD $k=2$



$\log |\Phi|$ vs. t/M at $x_q = 10M$ (RAD $k=2$)

- RAD PINN₂: τ error 1.80% vs. 1.50% (paper repro) — comparable accuracy.
- RAD primarily improves **spatial** accuracy (RMSD), not QNM extraction.

Zerilli $\ell = 2$ — QNM Extraction

	τ/M	ωM
TRUE	11.241	0.3737
FD₁	10.929 (2.78)	0.3716 (0.56)
FD₂	11.465 (2.00)	0.3798 (1.63)
PINN₁	10.831 (3.65)	0.3705 (0.85)
PINN₂	11.443 (1.80)	0.3790 (1.43)

FD_{1,2} identical to paper repro
(same FD solver).

PINN_{1,2} from RAD $k=2$ run.

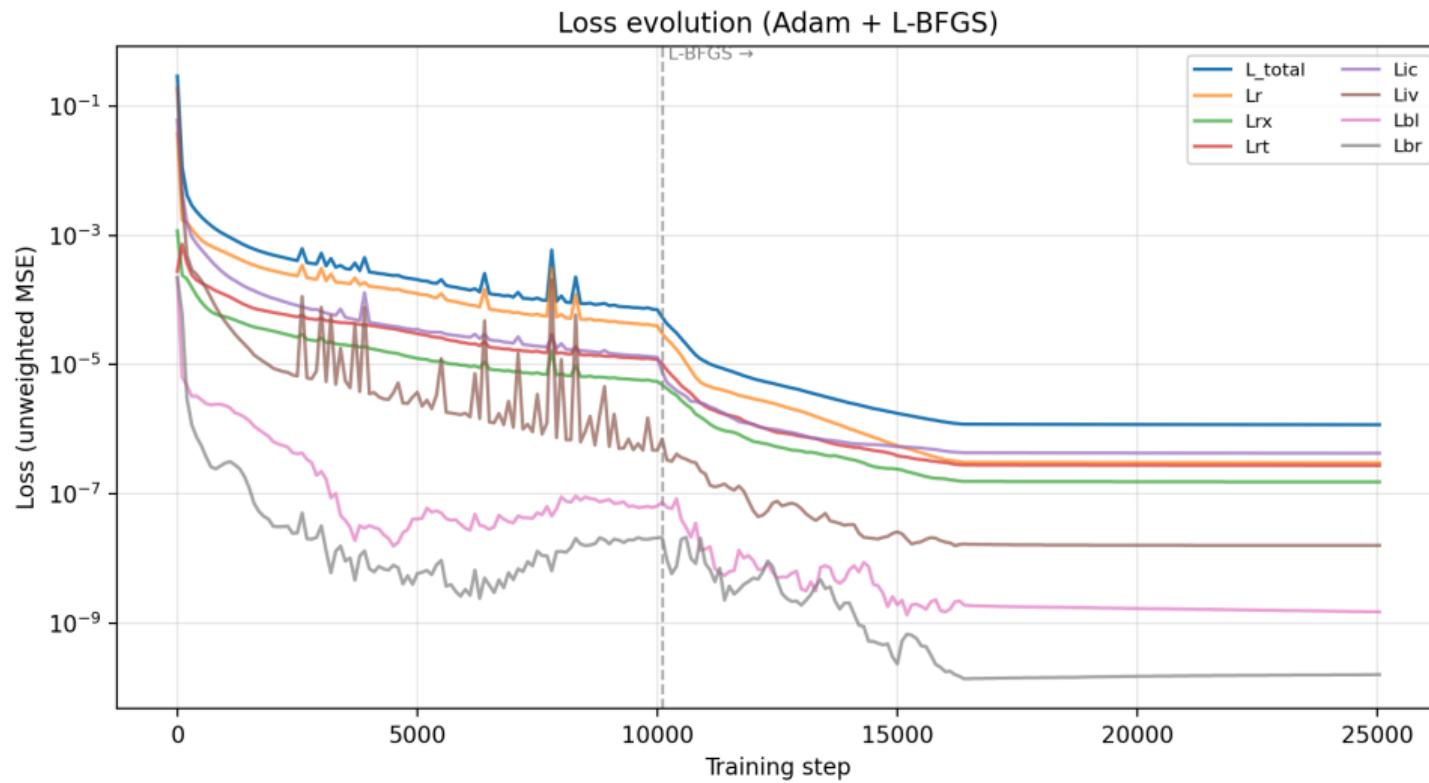
Summary of Findings

- ① **PINNs can compute QNMs** of Schwarzschild black holes to within $\sim 1\text{--}2\%$ of known values, using $15\times$ fewer collocation points than FD.
- ② **RAD adaptive sampling** improves solution accuracy by $\sim 37\%$ (RMSD) over uniform sampling, by concentrating points where the PDE residual is largest.
- ③ **Aggressive RAD** ($k = 2$) provides the best spatial accuracy (lowest RMSD/RL2). Anchor retention has the second-lowest loss but the highest RMSD among RAD runs — over-fitting to anchor points.
- ④ **Exponential reweighting fails** — naively amplifying late-time residuals destabilises training (τ error $> 20\%$).

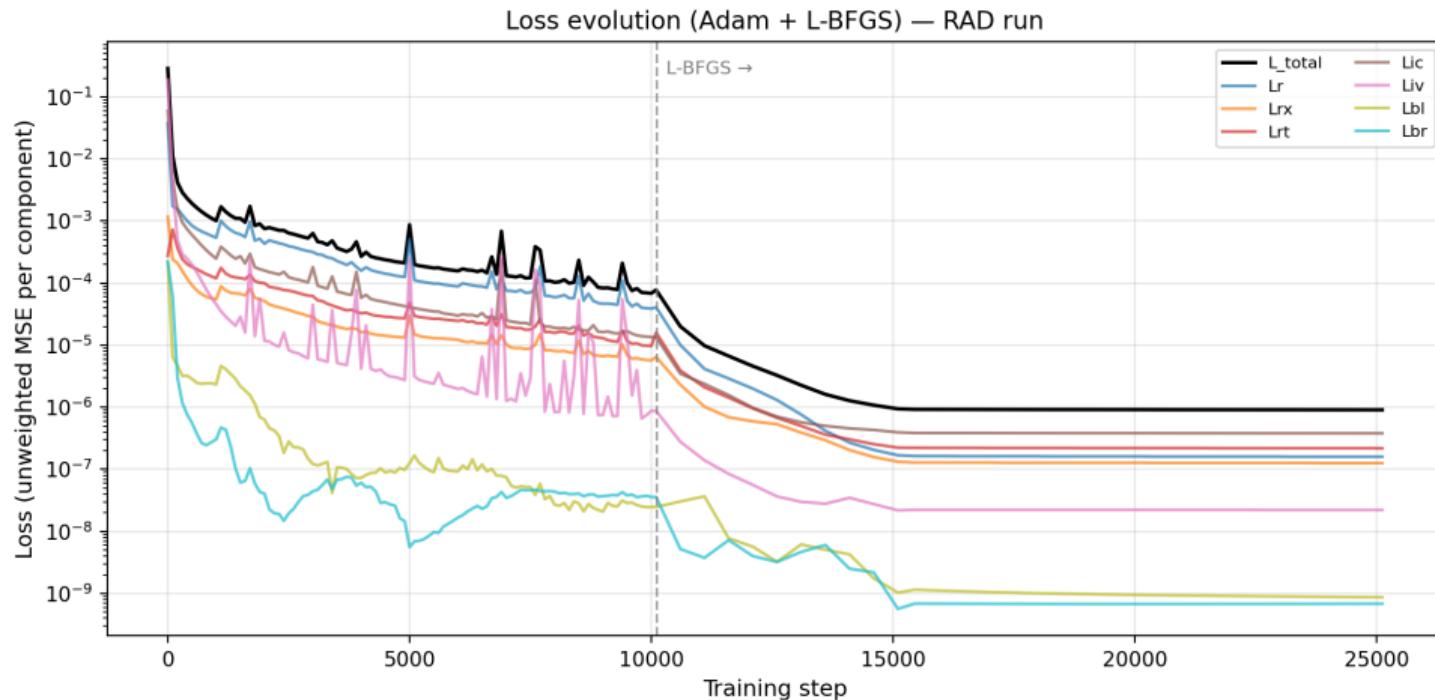
Thank you

Questions?

Appendix: Loss History — Reproduction Run (Uniform Sampling)



Appendix: Loss History — RAD Adaptive Sampling



Appendix: QNM Extraction — Method 1 (FFT + Envelope)

Extract ω and τ **independently** from the late-time waveform $\Phi(x_q, t)$.

Step 1 — Estimate ω via FFT:

- ① Window the signal between t_{start} and t_{end} ; subtract the mean.
- ② Apply a Hann window to reduce spectral leakage.
- ③ Zero-pad by $64\times$ and compute the real FFT.
- ④ Locate the magnitude peak; refine with **parabolic interpolation** for sub-bin accuracy.
- ⑤ Convert: $\omega = 2\pi f_{\text{peak}}$.

Step 2 — Estimate τ via log-linear envelope fit:

- ① Find all local maxima of $|\Phi(x_q, t)|$ (envelope peaks).
- ② Since the envelope decays as $A e^{-t/\tau}$, take the logarithm:

$$\ln |y_{\text{peaks}}| = \ln A - \frac{t}{\tau}.$$

- ③ A linear fit gives slope $m = -1/\tau$, hence $\tau = -1/m$.

Appendix: QNM Extraction — Method 2 (Nonlinear Curve Fit)

Fit the full damped-cosine model **simultaneously** to the waveform:

$$\Phi(x_q, t) \approx A e^{-t/\tau} \cos(\omega t + \phi).$$

Procedure:

- ① Obtain initial guesses from Method 1: ω_0, τ_0 .
- ② Set $A_0 = \max |\Phi|$, $\phi_0 = 0$.
- ③ Run `scipy.optimize.curve_fit` (Levenberg–Marquardt) with four free parameters (A, τ, ω, ϕ).

Advantages over Method 1:

- All four parameters are fitted **jointly**, so correlations between ω and τ are captured.
- Uses **every data point**, not just envelope peaks \Rightarrow typically more accurate.

Disadvantage:

- Requires good initial guesses; can fail to converge if the signal is noisy or the window is too short.

Appendix: L-BFGS — Motivation & Core Idea

Why two optimisers?

- **Adam** (Phase 1): stochastic, handles noisy/flat loss landscapes well; reaches a rough neighbourhood of the minimum quickly.
- **L-BFGS** (Phase 2): quasi-Newton method that uses curvature information to converge much faster once we are near a minimum.

Full BFGS recap. Newton's method updates parameters via $\theta_{k+1} = \theta_k - H_k^{-1} \nabla \mathcal{L}_k$, where H_k is the Hessian. Computing and inverting H is $O(n^2)$ storage and $O(n^3)$ per step (n = number of parameters) — **impractical for neural networks**.

BFGS (Broyden–Fletcher–Goldfarb–Shanno) avoids forming the Hessian explicitly. Instead it maintains an *approximate inverse Hessian* B_k^{-1} and updates it each step using only gradient differences:

$$s_k = \theta_{k+1} - \theta_k, \quad y_k = \nabla \mathcal{L}_{k+1} - \nabla \mathcal{L}_k.$$

Appendix: L-BFGS — The BFGS Update

The rank-2 update of the *approximate inverse Hessian* B_k^{-1} :

$$B_{k+1}^{-1} = (I - \rho_k s_k y_k^T) B_k^{-1} (I - \rho_k y_k s_k^T) + \rho_k s_k s_k^T, \quad \rho_k = \frac{1}{y_k^T s_k}.$$

This ensures B_{k+1} satisfies the **secant condition** $B_{k+1} s_k = y_k$ while staying symmetric positive-definite.

The search direction is then:

$$d_k = -B_k^{-1} \nabla \mathcal{L}_k$$

followed by a **line search** along d_k to find the step size.

Problem: Still $O(n^2)$ storage for the dense matrix B^{-1} — too expensive when n is the number of neural network weights.

Appendix: L-BFGS — Limited-Memory Variant

L-BFGS (Limited-memory BFGS) solves the storage problem:

- Store only the last m pairs $\{s_k, y_k\}$ (typically $m = 20\text{--}50$).
- Reconstruct the matrix–vector product $B_k^{-1}\nabla\mathcal{L}$ on-the-fly via a **two-loop recursion** — never form B^{-1} .
- Storage: $O(mn)$ instead of $O(n^2)$; each step is $O(mn)$.

Two-loop recursion (Nocedal 1980):

- ➊ *Backward pass*: for $i = k-1, \dots, k-m$, compute scalar $\alpha_i = \rho_i s_i^T q$ and set $q \leftarrow q - \alpha_i y_i$.
- ➋ *Scale*: $r \leftarrow H_k^0 q$ (diagonal initialisation).
- ➌ *Forward pass*: for $i = k-m, \dots, k-1$, set $r \leftarrow r + s_i(\alpha_i - \rho_i y_i^T r)$.
- ➍ *Search direction*: $d_k = -r$.

In our pipeline:

- Adam runs for 10 k steps \Rightarrow L-BFGS refines for 15 k steps.
- L-BFGS uses *full-batch* gradients (all 33 600 points), giving accurate curvature estimates.
- Typically reduces the loss by a further 1–2 orders of magnitude.

Appendix: Finite Difference Solver — Method of Lines

Rewrite the PDE $\Phi_{tt} - \Phi_{xx} + V\Phi = 0$ as a **first-order system**:

$$u = \Phi, \quad v = \Phi_t, \quad \begin{cases} u_t = v, \\ v_t = u_{xx} - V(x) u. \end{cases}$$

Spatial discretisation (2nd-order central differences):

$$u_{xx}|_i \approx \frac{u_{i+1} - 2u_i + u_{i-1}}{\Delta x^2}, \quad 1 \leq i \leq N_x - 2.$$

At boundaries ($i = 0$ and $i = N_x - 1$): 2nd-order **one-sided** stencils, e.g.

$$u_{xx}|_0 \approx \frac{2u_0 - 5u_1 + 4u_2 - u_3}{\Delta x^2}.$$

This converts the PDE into a system of ODEs in time: $\dot{\mathbf{U}}(t) = \mathbf{F}(\mathbf{U}(t))$, where $\mathbf{U} = (u_0, \dots, u_{N_x-1}, v_0, \dots, v_{N_x-1})$.

Appendix: Finite Difference Solver — RK4 & Boundary Conditions

Time integration — Classical RK4:

$$\mathbf{U}^{n+1} = \mathbf{U}^n + \frac{\Delta t}{6} (\mathbf{k}_1 + 2\mathbf{k}_2 + 2\mathbf{k}_3 + \mathbf{k}_4),$$

where $\mathbf{k}_1 = \mathbf{F}(\mathbf{U}^n)$, $\mathbf{k}_2 = \mathbf{F}(\mathbf{U}^n + \frac{\Delta t}{2}\mathbf{k}_1)$, etc.

4th-order accuracy in time; requires $\Delta t / \Delta x \leq 1$ (CFL condition). We use $\Delta x = 0.2M$, $\Delta t = 0.1M$
 \Rightarrow CFL ratio = 0.5.

Sommerfeld radiative BCs (applied at every RK4 sub-step):

$$\text{Left } (x \rightarrow -\infty) : (\partial_t - \partial_x)\Phi = 0 \Rightarrow v_0 = u_x|_0,$$

$$\text{Right } (x \rightarrow +\infty) : (\partial_t + \partial_x)\Phi = 0 \Rightarrow v_{N_x-1} = -u_x|_{N_x-1}.$$

Spatial derivatives u_x at boundaries use 2nd-order one-sided stencils:

$$u_x|_0 \approx \frac{-3u_0 + 4u_1 - u_2}{2\Delta x}, \quad u_x|_{N_x-1} \approx \frac{3u_{N_x-1} - 4u_{N_x-2} + u_{N_x-3}}{2\Delta x}.$$

Appendix: What Is the PINN Solution?

The PINN does **not** produce a closed-form expression. The “solution” is a neural network:

$$\Phi_{\theta}(x, t) = A \tanh\left(\mathbf{W}^{(5)} \tanh(\cdots \tanh(\mathbf{W}^{(1)} \begin{pmatrix} x \\ t \end{pmatrix} + \mathbf{b}^{(1)}) \cdots) + \mathbf{b}^{(5)}\right)$$

- 4 hidden layers: sizes [80, 40, 20, 10], tanh activation.
- Outer tanh bounds the output to $[-1, 1]$.
- $\theta = \{\mathbf{W}^{(l)}, \mathbf{b}^{(l)}\}$: **4,521 trainable parameters** (fixed after training).
- Evaluate at *any* (x, t) — no grid required.
- Differentiate *exactly* via automatic differentiation.

Key distinction:

- **FD solver**: returns Φ on a discrete grid $\{x_i, t_j\}$.
- **PINN**: returns a *continuous function* $\Phi_{\theta}(x, t)$ that can be queried at arbitrary resolution.
- Neither gives a human-readable formula — both are *numerical* solutions.